

CONFIDENTIAL · STRATEGIC — FOR DISCUSSION WITH QUAD / DOVER

Sovereign Harness vs. Agentic AI Governance Package

Alignment, Differences & Collaboration Opportunities

A comparison of **Decidr MCP** (decidrmcp.com) and the **Breathline Sovereign Harness** / Agentic AI Governance Package

Prepared by Kenn+ (Kenneth Mangum · KM-1176) · 2026-06-10 · Derived from sealed sources · Decidr details from public materials (to confirm with Dover)

Same conviction, two depths

Decidr MCP and the Sovereign Harness answer the same emerging need: **governed decision infrastructure** that keeps a human in command of AI agents — turning conversational generation into **reviewable, accountable decisions** with an audit trail. The alignment in philosophy is real and unusual: both treat **human judgment as the scarce resource to protect**, both make proposals **reviewable before commitment**, and both keep a record of what was decided.

They differ in **depth and scope**. Decidr is a lightweight, MCP-native commercial tool that bridges Claude conversations into governed initiatives and reviewable decisions — a fast, low-friction on-ramp. The Sovereign Harness is a full-stack sovereign substrate: receipted obligations on a hash-chained ledger, structural human gates (breath-gates), a cockpit of review surfaces (Atrium), a sovereign execution node, constitutional invariants enforced in the runtime, cryptographic anchoring, federation across nodes, and a living book↔code source of truth — designed for fiduciary, regulated, and multi-generational contexts.

The opportunity: these are complementary, not competing. A Decidr-class package is exactly the kind of human-friendly surface that could **license and layer on top of our sovereign primitives** — gaining verifiable governance, cryptographic auditability, and rigor it would otherwise have to build — while we keep the foundational research, the constitutional core, and the sovereign node as retained IP. This document maps where the two align, where ours goes deeper, and how a collaboration could be structured.

Already running — not a deck

The harness is proving value in live testbeds today: (1) a **regulated quality-systems consultancy** (welding coordination / QMS) — AI-assisted plan generation and gap analysis against real standards, an interactive natural-language edit loop with status tracking, and human engineering sign-off as the final gate, producing exportable signed artifacts with full provenance; and (2) a **collective economic coordination primitive** — privacy-preserving pooling of procurement intent, agent swarms surfacing optimized paths with full handoff traces, and an explicit human gate on any material commitment. Both demonstrate the stack handling real "should we proceed?" decisions with receipts. (Source: prepared Agentic AI Governance Package — Financial Case.)

2 · SIDE-BY-SIDE COMPARISON

Primitive-by-primitive

Dimension	Decidr MCP (decidrmcp.com)	Sovereign Harness / Agentic AI Governance Package
Decision / proposal unit	"Governed initiative" — an unstructured request converted into questions and reviewable decisions, kept connected to the work.	B32 obligation packet — a typed, receipted unit carrying level (L0–L3), role/capability, surfaces, tests, LGP-alignment, and exact source citations; double-entry (debit/credit) on a hash-chained ledger.
Review surface	Initiative planner, agent chat, message composing — reviewable decision documentation alongside the project.	Atrium cockpit — Kanban/board review surfaces, an "Awaiting-Me" human-gate queue, manuscript viewer, voice/text feedback capture, KDP dispatch, coherence/drift monitor — all rendering live node state.
Governance model	Governance "built in" — decisions stay connected to work with accountability (public framing).	Constitution@A1 + K1–K4 invariants enforced structurally in the runtime (Human-Primacy, Default-Deny, Audit-Immutable, Constitution-Validated-Extension); SOURCE/TRUTH/INTEGRITY triad.
Human gate	Human review of decisions before they are recorded/actioned.	Breath-gate — Propose → Approve → Execute; nothing material commits without a documented human authorization token. One recurring gate by design.
Auditability	Audit trails / decision accountability (per public materials).	Receipted, hash-chained ledger + cylinder seal chain; evidence tiers (E0 rejected · E1 pointer · E2 artifact+hash); Merkle anchoring; independently replayable.
Source of truth	Decision documents connected to project work.	Living book ↔ code — the published books are the spec; code is their derived form; surfaces render from sealed sources (render, not recreate), with a fidelity gate proving the trace.
Alignment objective	Better, governed project decisions; accountability.	Lasting Generational Prosperity (LGP) — every packet, gate, and surface scored by its LGP / human-primacy contribution; multi-generational continuity is a design goal.
Architecture & AI plumbing	MCP-native CLI integrating Claude via Model Context Protocol; cloud product.	Sovereign node (loopback-trust, local-first), a Node API contract seam, sovereign inference routing by sensitivity; replaceable cockpit over a sovereign contract.
Scope	Focused commercial package — project/initiative decision governance for teams; quick to adopt.	Full sovereign stack — governance, execution, economics (yield organism), federation, and a productizable commercial layer on top.

Dimension	Decidr MCP (decidrmcp.com)	Sovereign Harness / Agentic AI Governance Package
Strengths	Low friction, fast on-ramp, MCP-standard, commercially packaged and available today; clean product surface.	Depth, verifiability, cryptographic auditability, constitutional enforcement, rigor standards, sovereignty & portability, generational durability.
Gaps (candid)	Depth of audit/cryptographic proof, sovereignty/portability, structural (not policy) governance, and rigor against "correct-but-hollow" output are unproven from public materials.	Heavier to adopt; commercial packaging/GTM still maturing; the lightweight, instantly-usable on-ramp is exactly where a Decidr-class product is ahead.

Decidr column reflects public materials at decidrmcp.com (the page is intentionally sparse); specifics to confirm directly with Dover. Sovereign Harness column derives from sealed specs (Review-Ready Rail, Board Rigor Standard, Controlled-Link & Book↔ Code Tree, Constitution@A1) and the live node.

Where the concepts match

- **Reviewable proposals before commitment.** Decidr's "reviewable decisions" map directly onto our **obligation packets + breath-gate** — a proposal is examined and dispositioned before it becomes action.
- **Human primacy.** Both insist a human ratifies; agents inform, humans decide. Our K1 makes this a structural runtime guarantee rather than a workflow convention.
- **From unstructured input to governed unit.** Decidr "a few sentences → a governed initiative" is the same move as our **capture** → **packet** → **validate** → **gate** governed loop.
- **Audit trail.** Both keep a record of decisions; ours is hash-chained, receipted, and independently replayable.
- **MCP common ground.** Decidr is MCP-native; our node exposes a clean API contract and can speak MCP — a natural technical bridge.

The differences, along five axes

Axis	Decidr MCP	Sovereign Harness
Foundation	A product feature on the MCP/Claude stack — governance as application logic atop a cloud tool.	A sovereign substrate — Constitution@A1 + a runtime that enforces governance structurally , local-first and portable.
Source of truth	Decision documents connected to project work.	Living book↔code — published books are the spec; surfaces render from sealed text; a fidelity gate proves the trace.
Sovereignty depth	Cloud SaaS; data and governance live in the vendor's product.	Operator owns the node, data, governance, and audit chain; "contract sovereign, UI replaceable."
Generational continuity	Project/initiative horizon.	Designed to be inherited across stewards; scored against Lasting Generational Prosperity.
Federation	Single-tenant product workflow (per public materials).	Sovereign nodes coordinate by signed resonance , not central control — multi-party without merging chains.

Where ours goes deeper

- **Living book ↔ code source.** Our governance is derived from published books that are the spec; a Book ↔ Code tree audits that every module traces to a chapter, and surfaces render from sealed text — provable, not asserted.
- **Helix rendering & derived surfaces.** Operator docs (seeit) and review surfaces are *rendered* from sealed sources, auto-refreshing — no parallel canon to drift.
- **Merkle anchors & receipted memory.** A tamper-evident, hash-chained record with cryptographic anchoring; you can prove what happened years later.
- **K1–K4 constitutional invariants.** Human-Primacy, Default-Deny, Audit-Immutable, Constitution-Validated-Extension — enforced server-side as structural properties, defensible to a board or regulator in ~60 seconds on the operator's own node.
- **Federation resonance.** Sovereign nodes coordinate by signed resonance, not central control — multi-party governance without merging chains.
- **Generational continuity.** The whole stack is scored against Lasting Generational Prosperity and built to be inherited across stewards.
- **Rigor standards (no rubber stamps).** A Review-Ready Rail with machine-checked board rigor (LGP alignment per finding, findings → obligations, depth gate, fidelity witness, human-sense test) — the antidote to "correct-but-hollow" output.

How a Decidr-class package could layer on our primitives

The cleanest model: the commercial package keeps its fast, friendly surface and **licenses the sovereign substrate underneath** — gaining verifiable governance and auditability it would otherwise build from scratch, while our core IP (constitutional design, sovereign node source, book↔code research) stays retained.

Technical integration points

- **B32 packets + obligation ledger:** their initiatives mint **B32 obligation packets** on our hash-chained ledger — typed, receipted, evidence-tiered decision records.
- **MCP seam:** expose our Node API + ledger as an MCP server; Decidr-style agents emit proposals straight into our governed loop.
- **Atrium cockpit:** white-label or co-branded review surfaces (Awaiting-Me gate, board review, coherence monitor) over their initiatives.
- **Helix rendering:** their decision docs render from sealed sources (render-not-recreate) — operator-facing surfaces that can't drift from the source.
- **Merkle enforcement:** decisions anchor to our Merkle/hash-chain so each record is tamper-evident and independently replayable.

Value flow (backer / royalty)

- **Licensing/royalty** on the sovereign substrate embedded in the commercial package.
- **Equity or revenue share** in a separate commercial vehicle (the Agentic AI Governance Package entity) — preferred for aligned capital.
- **Performance participation** tied to governed "proceed" outcomes (fewer false starts, faster ratified projects, auditable sign-off).
- **Future on-chain yield instruments** linked to the governance packages, once regulatory clarity exists.

Structure mirrors the prepared Agentic AI Governance Package financial case: core retained in the primary entity; production layer licensed to the commercial vehicle.

1

Pilot bridge

2

Co-branded surface

3

Commercial vehicle

Stand up the MCP seam; route one Decidr-style initiative through our obligation ledger + breath-gate; produce a receipted decision.

Wrap it in an Atrium review surface (Awaiting-Me + board review) the partner can demo to a regulated client.

Formalize licensing + value flow; the package sells, the sovereign core stays retained and continues its research cadence.

5 · RISKS & RECOMMENDATIONS

Where shallow work creeps in — and how our rigor strengthens the offering

Risk — "correct but hollow." A lightweight decision bridge can ship outputs that *look* governed (a record exists, a human clicked approve) but are shallow — no depth check, no source trace, no proof the decision was substantive. At scale this compounds into GIGO and erodes trust exactly where accountability was the selling point.

Recommendation. Our **Board Rigor Standard** and **validation/fidelity harness** are the structural antidote — and a differentiator a partner could license: every material finding becomes a tracked obligation, every claim must trace to a resolving source, a depth gate forbids silent "all clear," and an independent witness verifies before anything reaches the human. This is what turns "an audit trail" into "verifiable governance."

IP & sovereignty protection. Retain the sovereign node source, constitutional design, and book↔code research in the primary entity; license only the production surfaces, templates, and coordination modules. The Controlled-Link Doctrine (published artifacts point only at surfaces we control) keeps the partner's deployments dependent on our living, governed source — not a fork.

Next steps for the Dover / Quad discussion

- Confirm Decidr's actual depth (audit model, governance enforcement, MCP scope) directly with Dover — this doc's "their" column is from public materials.
- Agree the comparison framing: **complementary layers**, not competitors — Decidr (or their package) as a friendly surface; our harness as the verifiable substrate.
- Scope a **pilot bridge** (Step 1 above): one initiative through the MCP seam → receipted, breath-gated decision → demoed in an Atrium surface.
- Outline the commercial-vehicle + value-flow terms (licensing/royalty/equity) consistent with the prepared financial case.

Confidential — prepared for KM-1176 for discussion with Quad / Dover. Sovereign Harness claims derive from sealed sources (Constitution@A1; Review-Ready Rail, Board Rigor, Controlled-Link & Book↔Code Tree specs; the live node). Decidr MCP claims derive from public materials at decidrmcp.com and are marked for confirmation with Dover. Framing remains consistent with human primacy and Lasting Generational Prosperity. ·

∞Δ∞